




Red-Black Trees

Comprehensive Learning Module

A deep dive into one of computer science's most elegant and widely-deployed data structures — powering the Linux kernel, Java collections, and C++ STL.

ADVANCED

DATA STRUCTURES

ALGORITHMS

Topic Classification

Field	Details
Topic Name	Red-Black Trees
Category	Data Structures
Domain	Computer Science / Self-Balancing Binary Search Trees
Topic Type	Conceptual + Algorithmic + Implementation
Difficulty Level	Advanced
Prerequisites	Binary Trees, BST, AVL Trees, Recursion
Industry Relevance	Very High — Linux kernel, Java TreeMap, C++ std::map

Recommended Learning Path



BST

Binary Search Trees — the foundation



AVL Trees

First self-balancing BST



Red-Black Trees

You are here



B-Trees

Disk-based multiway trees

What Is a Red-Black Tree?

Formal Definition

A Red-Black Tree (RBT) is a self-balancing binary search tree in which every node carries an additional **color bit** (RED or BLACK). A set of **five structural invariants** must hold at all times, collectively ensuring the tree height stays within $2 \times \log_2(n+1)$ — guaranteeing $O(\log n)$ operations.

Simple Explanation

Imagine a BST where every node wears either a red or black hat. The rules about which hats can go where are carefully chosen so that the tree can never grow too lopsided — yet the rules are relaxed enough that fixing a violation usually requires fewer reorganization steps than AVL trees. It's the balance between "strict enough to be fast" and "flexible enough to be efficient on writes."

Core Purpose


- Guarantee $O(\log n)$ worst-case for search, insert, and delete
- Fewer rotations per insert/delete than AVL trees → preferred for write-heavy workloads
- Used as the standard map/set implementation in most production languages and OS kernels

Key Terminology

Term	Definition
Color	Each node is RED or BLACK
Black Height (bh)	Number of black nodes on any path from a node to a NIL leaf
NIL Sentinel	A shared dummy leaf node, always BLACK
Rotation	Left or right restructuring (same as AVL)
Recoloring	Changing a node's color to restore properties
Uncle Node	Sibling of the parent of the newly inserted node
Double Black	Temporary state after deletion requiring special handling

i The five Red-Black Properties together guarantee: $\text{height}(\text{tree}) \leq 2 \times \log_2(n+1)$

The Five Red-Black Properties



Why Does It Exist?

Problem 1 — AVL is Too Strict for Write-Heavy Workloads

AVL trees can trigger $O(\log n)$ rotations on deletion. Red-Black Trees need at most **3 rotations per insert** and **3 rotations per delete**, drastically reducing restructuring overhead.

Problem 2 — Production Systems Need a Write-Balanced Tree

OS kernels, language runtimes, and databases perform many more insertions and deletions than pure lookups on their internal ordered maps. Red-Black Trees provide the right write-performance balance.

Problem 3 — Simpler Recoloring vs. Rotation

Recoloring (flipping a node's color) is $O(1)$ and requires no pointer changes. Many Red-Black fix-ups are pure recolorings, making them cheaper than AVL height-update + rotation chains.

Real-World Motivation

Linux CFS Scheduler

Processes ordered by virtual runtime in a Red-Black Tree — runs on every Linux machine globally

Java TreeMap/TreeSet

Explicitly documented as using Red-Black Trees in the Java standard library


C++ STL `std::map/set`

Red-Black Tree in every major implementation — GCC libstdc++, Clang libc++

Nginx

Uses Red-Black Trees for timer management in its event loop

Evolution & History



Connection to B-Trees

A Red-Black Tree is isomorphic to a 2-3-4 B-Tree. This mathematical equivalence is why Red-Black Trees have the same logarithmic depth guarantees as B-Trees.

- A BLACK node with two RED children = a 4-node in the B-Tree
- A BLACK node with one RED child = a 3-node
- A BLACK leaf = a 2-node


Variants

Variant	Key Difference
Classic RBT	Both left and right red links allowed
Left-Leaning RBT	Red links only go left; half the cases to handle
AA Tree	Simpler rules; only right red links allowed
Concurrent RBT	Fine-grained locking or lock-free variants

How It Works

Internal Node Structure

Node:



NIL sentinel: shared BLACK leaf
(all NULL pointers point here)

Insertion Process

01

BST Insert

Insert like a standard binary search tree

02

Color RED

Color the new node RED

03

Fix-Up

Fix violations of Property 2 and Property 4

04

Root BLACK

Recolor root to BLACK (Property 2 guarantee)

Black Height Verification

$$\text{bh}(\text{node}) = \text{bh}(\text{node.left}) + (1 \text{ if node.left is BLACK else } 0)$$

All paths root $\rightarrow$ NIL must have equal black height

Insertion Fix-Up — Three Cases

Case 1: Uncle is RED (Recoloring)

Recolor parent and uncle to BLACK, grandparent to RED. Move problem up to grandparent.



Case 2: Uncle is BLACK, Z is Inner Child (LR or RL)

Rotate parent in opposite direction to convert to Case 3.



Case 3: Uncle is BLACK, Z is Outer Child (LL or RR)

Rotate grandparent and recolor. This resolves the violation.

Deletion Fix-Up Cases

DB = Double Black node, S = Sibling

Case 1: S is RED $\rightarrow$ Rotate + recolor $\rightarrow$ converts to Case 2/3/4

Case 2: S is BLACK, S's children BLACK $\rightarrow$ Recolor S RED,
push DB up

Case 3: S is BLACK, far child BLACK, near child RED →
Rotate + recolor → Case 4

Case 4: S is BLACK, far child RED → Rotate grandparent, recolor → Done

Types & Variants

Variant	Balance Guarantee	Rotations (Insert)	Rotations (Delete)	Best For
Classic RBT	$\text{height} \leq 2\log(n+1)$	≤ 2	≤ 3	General purpose
LLRB (Left-Leaning)	Same	≤ 2	≤ 2	Simpler implementation
AA Tree	Slightly relaxed	≤ 2	≤ 2	Functional languages
Concurrent RBT	Same	≤ 2	≤ 3	Multithreaded systems



Classic RBT

The standard formulation allowing red links in both directions. Most widely documented and implemented in production systems like Linux and Java.



Left-Leaning RBT

Introduced by Sedgwick in 1987. Red links only go left, cutting the number of cases to handle roughly in half. Ideal for teaching and simpler codebases.



Concurrent RBT

Fine-grained locking or lock-free variants designed for multithreaded environments where multiple threads may read and write simultaneously.

Use Case 1: Linux Kernel CFS Scheduler

The Problem

The Linux kernel must schedule thousands of processes and threads fairly, selecting the next process to run in **$O(\log n)$** time with **$O(1)$** minimum-finding.

The Solution

Each runnable process has a "virtual runtime" key stored in a Red-Black Tree, ordered by virtual runtime. The scheduler always picks the leftmost node — the process with the smallest virtual runtime.

Why RBT?

- Frequent inserts (new processes arriving)
- Frequent deletes (processes finishing)
- Always picks leftmost node (minimum virtual runtime = next to run)
- $O(\log n)$ insert/delete and $O(1)$ leftmost node access (cached)

Architecture

Run Queue **Per** CPU:

RB Tree keyed **by** vruntime

└── leftmost cached ($O(1)$ access)

└── **insert/delete** $O(\log n)$

✅ **Impact:** Linux CFS scales from embedded systems to 1000-core servers using the same data structure. This is the scheduler running inside every Android phone, AWS EC2 instance, and Linux server globally.

Key Insight

The Linux kernel packs the color bit into the lowest bit of the parent pointer — since pointers are word-aligned, the lowest bit is always 0 for valid addresses. This means **color is stored for free** with no extra memory per node.

Use Case 2: Java TreeMap / TreeSet

The Problem

Provide a sorted map and sorted set with guaranteed $O(\log n)$ operations, supporting `headMap()`, `tailMap()`, `floorKey()`, `ceilingKey()` range operations across millions of Java applications.


Why RBT over AVL?

Java's workload involves heavy concurrent mutations (wrapped via `Collections.synchronizedMap`). Fewer rotations per mutation reduces lock contention time — a critical advantage in enterprise environments.

Benefits

- Used in millions of Java applications
- Powers event schedulers, trading systems, concurrent priority queues
- All `NavigableMap` API operations run in $O(\log n)$

Architecture Stack



- 1 **`java.util.TreeMap`**
- 2 **`Entry<K,V>[]` + RBT structure**
- 3 **Red-Black fix-up on `put()/remove()`**
- 4 **`NavigableMap` API: `floorKey`, `ceilingKey`, `subMap` — all $O(\log n)$**

Use Case 3: C++ STL

std::map Internals

The Problem

C++ programs need a sorted associative container with predictable worst-case performance across systems programming, game development, and financial applications.

The Solution

GCC's libstdc++ and Clang's libc++ both implement `std::map` and `std::set` with Red-Black Trees. Every C++ developer implicitly uses Red-Black Trees whenever they use these containers.

Affected Containers

`std::map`

`std::set`

`std::multimap`

`std::multiset`

Internal Structure

```
std::map<Key, Value>
→ _Rb_tree<Key, Value,
Compare, Allocator>
→ RBT node: {
    _M_color,
    _M_parent,
    _M_left,
    _M_right,
    _M_value
}
```

i Consistent $O(\log n)$ performance across all STL sorted containers — guaranteed by the Red-Black Tree invariants baked into every major C++ standard library implementation.

Use Case 4: Java HashMap Treeification

The Problem


Java's HashMap uses linked list buckets for collision handling. Under hash collisions, one bucket can grow to $O(n)$ lookup — a serious performance and security vulnerability.

The Solution

Since **Java 8**, when a single bucket exceeds **8 entries**, it converts the linked list into a Red-Black Tree, reducing worst-case lookup in that bucket from $O(n)$ to $O(\log n)$.

Why RBT?


The treeification and untreeification (back to list when entries drop below 6) happens frequently. RBT's efficient insert/delete is ideal for this dynamic conversion pattern.



⚠ Security Benefit: Prevents hash-collision denial-of-service attacks and performance cliffs in HashMap-heavy applications. Adversaries can no longer craft inputs that degrade HashMap to $O(n)$.

mHam Bucket Colles

eat Vi Iruscunaiiaiom,



alahocl lto araeY

Implementation 1: Node Structure & Color Constants

Pseudocode

```
CONSTANT RED = 0
CONSTANT BLACK = 1

CLASS RBNode:
    key
    color = RED    // New nodes always start RED
    left = NIL
    right = NIL
    parent = NIL

NIL_SENTINEL = new RBNode()
NIL_SENTINEL.color = BLACK // NIL is always BLACK
```

Design Decisions

New nodes start RED

Inserting as RED only risks violating Property 4 (consecutive reds), which is easier to fix than a black-height violation from inserting BLACK.

Shared NIL Sentinel

A single shared BLACK sentinel node eliminates null checks throughout the fix-up code, simplifying all edge cases.

Parent Pointer

Required for iterative fix-up traversal back toward the root. Adds one pointer per node but enables $O(\log n)$ iterative operations.

Implementation 2: Left and Right Rotations


Left Rotate

```
FUNCTION left_rotate(tree, x):
  y = x.right
  x.right = y.left
  IF y.left != NIL:
    y.left.parent = x
  y.parent = x.parent
  IF x.parent == NIL:
    tree.root = y
  ELSE IF x == x.parent.left:
    x.parent.left = y
  ELSE:
    x.parent.right = y
  y.left = x
  x.parent = y
```

Right Rotate

```
FUNCTION right_rotate(tree, x):
  y = x.left
  x.left = y.right
  IF y.right != NIL:
    y.right.parent = x
  y.parent = x.parent
  IF x.parent == NIL:
    tree.root = y
  ELSE IF x == x.parent.right:
    x.parent.right = y
  ELSE:
    x.parent.left = y
  y.right = x
  x.parent = y
```

❏ Rotations in RBT are identical to AVL rotations — **left_rotate** and **right_rotate**. In RBT they are accompanied by color swaps. During rotation, no color changes occur — colors follow the nodes.



Implementation 3: Insert with Fix-Up

Insert Function

```
FUNCTION insert(tree, key):
  z = new RBNode(key)
  z.left = z.right = z.parent = NIL

  // Standard BST insert
  y = NIL
  x = tree.root
  WHILE x != NIL:
    y = x
    IF z.key < x.key: x = x.left
    ELSE:           x = x.right
  z.parent = y

  IF y == NIL:      tree.root = z
  ELSE IF z.key < y.key: y.left = z
  ELSE:            y.right = z

  // Fix Red-Black properties
  insert_fixup(tree, z)
```

Insert Fix-Up

```
FUNCTION insert_fixup(tree, z):
  WHILE z.parent.color == RED:
    IF z.parent == z.parent.parent.left:
      y = z.parent.parent.right // Uncle
      IF y.color == RED:         // Case 1
        z.parent.color = BLACK
        y.color = BLACK
        z.parent.parent.color = RED
        z = z.parent.parent
      ELSE:
        IF z == z.parent.right: // Case 2
          z = z.parent
          left_rotate(tree, z)
        z.parent.color = BLACK // Case 3
        z.parent.parent.color = RED
        right_rotate(tree, z.parent.parent)
    ELSE:
      // SYMMETRIC (right side mirror)
      // ... mirror image of above
      tree.root.color = BLACK // Property 2
```

Walkthrough — Insert [10, 20, 30]

01

Insert 10

Tree: 10(B) — colored BLACK immediately (root rule, Property 2)

02

Insert 20

BST: 10(B) → right → 20(R). Fix-up: 20's parent (10) is BLACK → no violation → done.

03

Insert 30 — Violation Detected

BST: 30(R) as right child of 20(R). Violation: consecutive reds 20(R) → 30(R). Uncle = NIL(B), grandparent = 10(B). Uncle is BLACK, z is outer child (RR Case → Case 3).

04

Fix-Up Applied

Recolor: 20 → BLACK, 10 → RED. Left rotate on 10. Result: 20(B) with children 10(R) and 30(R). All 5 properties hold ✓

After rotation:

```

20(B)
 /  \
10(R) 30(R)

```

Root recolored BLACK → already BLACK ✓

All 5 properties hold ✓

Hands-On Lab 1: Insert & Property Verification

BEGINNER-INTERMEDIATE

Objective

Implement RBT insert and verify all 5 Red-Black properties after each insertion.

Steps

- 1

Implement RBNode
Fields: key, color, left, right, parent
- 2

Create NIL Sentinel
Shared BLACK sentinel node
- 3

Implement Rotations
left_rotate and right_rotate with parent-pointer updates
- 4

Implement insert + insert_fixup
Full insertion with fix-up logic
- 5

Implement verify_properties(root)
Check all 5 properties recursively
- 6

Insert Sequence
Insert [7, 3, 18, 10, 22, 8, 11, 26] and verify after each

Expected Final State

Level 0: 10(B)
Level 1: 7(R) 18(R)
Level 2: 3(B) 8(B) 11(B) 22(B)
Level 3: 26(R)
Black height = 2 for all paths ✓

Troubleshooting

- ⚠

Property 4 violation → fix-up not running or wrong case selected
- ⚠

Property 5 violation → rotation not updating parent pointers correctly
- ℹ

NIL nodes must all point to the shared BLACK sentinel, not actual NULL

Learning Outcome

Deep understanding of when each fix-up case triggers and how the tree self-balances after each insertion.

Hands-On Lab 2: Delete & Edge Cases

ADVANCED

Objective

Implement full RBT deletion with Double Black fix-up.

Prerequisites

Complete Lab 1; understand the transplant operation.

Steps

01

Implement `transplant(tree, u, v)`

Replaces subtree rooted at `u` with subtree rooted at `v`

03

Implement `delete_fixup(tree, x)`

Handle all 4 sibling cases for Double Black resolution

Validation Checkpoints

Learning Outcome

Grasp why delete is significantly more complex than insert in Red-Black Trees, and how the Double Black abstraction simplifies the fix-up logic by reducing it to 4 well-defined cases.

The Double Black concept is a temporary bookkeeping state — it never persists in the final tree. It simply tracks that a black unit needs to be "absorbed" somewhere in the tree.

02

Implement `delete(tree, key)`

Find node, handle two-children case with inorder successor, track original color

04

Build and Delete

Build from [20, 10, 30, 5, 15, 25, 35], then delete 10, 30, 20

After deleting 10

Tree should reorganize; verify black height = 2

After deleting 30

Verify no double-black condition remains in the tree

After deleting 20

New root must be BLACK (Property 2)

Advantages & Disadvantages

Advantages

Perspective	Advantage
Performance	$O(\log n)$ guaranteed for all operations
Write Efficiency	At most 2 rotations for insert; at most 3 for delete
Industry Standard	Proven in Linux kernel, Java, C++ STL for decades
Flexibility	Balanced enough for both read and write workloads
Augmentability	Easily extended (interval trees, order-statistic trees)

Disadvantages

Perspective	Disadvantage
Complexity	Deletion fix-up has 6 cases (3 symmetric pairs); extremely difficult to implement correctly from scratch
Lookup Speed	Relaxed balance means slightly taller trees → marginally slower reads vs. AVL
Memory	Requires parent pointers in most implementations (extra pointer per node)
Debugging	Subtle violations in color assignment are hard to diagnose without a full property checker

Performance Considerations

Complexity Analysis

Operation	Time Complexity	Rotations
Search	$O(\log n)$	0
Insert	$O(\log n)$	≤ 2
Delete	$O(\log n)$	≤ 3
Space	$O(n)$	—

Height Guarantee

$\text{height(RBT)} \leq 2 \times \log_2(n + 1)$

For $n = 1,000,000$ nodes:

AVL **max** height ≈ 29

RBT **max** height ≈ 40

In practice (**random** data):

RBT $\approx 1.2 \times$ AVL height

Insert/Delete Rotations vs. AVL

Scenario	AVL Rotations	RBT Rotations
Best case insert	0	0
Worst case insert	2	2
Best case delete	0	0
Worst case delete	$O(\log n)$	3

✔ Red-Black Trees have **bounded** delete rotations (≤ 3); AVL does not. This is the key advantage for write-heavy workloads.

Benchmarking Notes

- Read-heavy (99% read):** AVL outperforms RBT by 5–15% due to shorter height
- Mixed workloads (50/50):** RBT matches or outperforms AVL
- Write-heavy (99% write):** RBT outperforms AVL significantly

Scalability & Reliability

Scale Reference

Scale	Nodes	RBT Height	Notes
Small	≤ 10K	≤ 28	Trivially fast
Medium	10K–10M	≤ 48	Excellent
Large	10M–1B	≤ 62	Still in-memory only
Massive	> 1B	—	Shard across instances; or use B+ Trees

High Availability Patterns

Persistent RBT

Path-copying on modifications preserves old versions — $O(\log n)$ space per modification. Used in Haskell and Clojure for persistent sorted maps.

Read-Write Lock

Multiple concurrent readers; single writer with RBT lock. Common pattern in kernel use for shared data structures.

Copy-on-Write per Epoch

Writer builds a new tree version; atomically swaps root pointer when done. Readers see a consistent snapshot throughout their operation.

Design & Architectural Considerations

Key Architecture Decisions

Decision	Trade-off
NIL Sentinel vs. NULL	Sentinel eliminates null checks in fix-up; adds 1 dummy node
Parent Pointers	Required for iterative fix-up; adds pointer per node
Recursive vs. Iterative	Iterative is production-safe for large trees; recursive simpler to understand
Color as Bit vs. Boolean	1-bit color can be packed into lowest bit of a pointer (used in Linux kernel)
Classic vs. LLRB	LLRB cuts case count in half; classic is more widely documented

Linux Kernel Optimization

The Linux kernel packs the color bit into the lowest bit of the parent pointer (since pointers are word-aligned, the lowest bit is always 0 for valid addresses):

```
parent_color = (parent_pointer | color_bit)
→ Color stored "for free" — no extra memory per node
```

Relevant Design Patterns

Interval Tree

Augment each node with max_end of subtree → $O(\log n)$ interval overlap queries (used in memory management)

Order-Statistic Tree

Augment with subtree_size → $O(\log n)$ rank/select operations

Event-Driven Timer Heap

Red-Black Trees as timer heaps in event loops (Nginx)

Security Considerations

Threat	Impact	Mitigation
Algorithmic Complexity Attack	Adversary crafts inputs to maximize rotations	RBT guarantees ≤ 3 rotations regardless of input $\rightarrow$ immune
Hash Collision $\rightarrow O(n)$ HashMap	Degraded Java HashMap lookup	Java 8+ treeifies bucket $\rightarrow$ RBT prevents $O(n)$ worst case
Memory Exhaustion	Unbounded inserts exhaust RAM	Enforce max node count; evict on capacity
Pointer Corruption	Corrupted parent/child pointers break invariants silently	Validate RB properties at runtime in debug builds
Color Bit Tampering	Low-level buffer overflows flip color bits	Full property verification on deserialization

 **Key Security Advantage:** Unlike hash tables, Red-Black Trees provide **worst-case $O(\log n)$** guarantees regardless of input distribution. This makes them inherently resistant to algorithmic complexity attacks that exploit hash collisions or degenerate BST inputs.

Best Practices

✓ Do's

→ Use a NIL sentinel node

Avoids null checks in fix-up cases and simplifies all edge-case handling

→ Maintain parent pointers

Required for iterative implementation at scale; enables $O(\log n)$ fix-up traversal

→ Implement `verify_rb_properties()`

Run it in all test cases — essential for catching subtle bugs

→ Write symmetric cases explicitly

Don't try to unify left/right with a variable — it reduces clarity and introduces bugs

→ Cache the minimum node pointer

$O(1)$ minimum access — the Linux CFS trick for leftmost node

→ Emit black height as a metric

In production, sudden changes indicate tree corruption

✗ Don'ts

→ Don't insert with RED root

Always force root to BLACK after fixup — Property 2 must always hold

→ Don't skip the NIL sentinel

Handling actual NULL requires extra null checks everywhere in fix-up code

→ Don't use RBT for disk storage

B+ Trees are the right choice for disk-based ordered storage

→ Don't debug deletion by hand

For more than 3 nodes, always write property verification code instead

→ Don't share RBT between threads without synchronization

Color changes are not atomic — concurrent access causes silent corruption

Common Mistakes & Pitfalls

Beginner Mistakes

Mistake	Why It Happens	Impact	Fix
New node colored BLACK	Intuitive but wrong	Violates black height immediately on insert	Always insert new nodes as RED; fix-up handles it
Null checks instead of NIL sentinel	Trying to avoid the sentinel	Fix-up code becomes bloated with null checks	Use a shared NIL sentinel node colored BLACK
Forgetting to color root BLACK after fix-up	Fix-up loop exits; root stays RED	Property 2 violated	Always set root.color = BLACK after fix-up


Intermediate Mistakes

Mistake	Why It Happens	Impact	Fix
Not updating parent pointers in rotation	Forgetting y.parent = x.parent step	Traversal to root breaks; fix-up loops infinitely	Carefully follow rotation algorithm; test with parent-walk after every insert
Confusing uncle lookup in symmetric case	The symmetric case mirrors left-side	Fix-up applies wrong rotation direction	Write both sides explicitly; don't try to generalize
Missing Case 2 in deletion	Case 2 is easy to skip; seems harmless	Double-black propagates to root; tree stays invalid	Implement all 4 deletion fix-up cases


Advanced / Production Mistakes

Mistake	Why It Happens	Impact	Fix
Using recursive fix-up at scale	Recursion is conceptually simpler	Stack overflow for n > 100K	Iterative fix-up with parent pointers
Not serializing color on persistence	Color treated as runtime state only	Deserialized tree has all nodes BLACK → not balanced	Serialize color bit with each node; verify properties on load
Race condition on concurrent access	RBT color changes look harmless	Silent corruption of tree structure	Use read-write lock or a concurrent-safe wrapper


Comparison with Alternatives




Red-Black Tree: Write-Write-balanced in-mer. $O(\log n)$ all ops. ≤ 3 delete rotations.




AVL Tree: Read-heavy in-memory. $O(\log n)$ all ops. $O(\log n)$ delete rotations.




B+ Tree: Disk-based storage. $O(\log n)$ all ops. Best for range queries.



Skip List: Concurrent systems. $O(\log n)$ expected. No rotations.



Hash Table: Unordered lookup. $O(1)$ average search. No ordering.



Feature Matrix

Feature	RBT	AVL	B+ Tree	Hash Table
Height Guarantee	$\leq 2\log(n)$	$\leq 1.44\log(n)$	$\leq \log_t(n)$	N/A
Delete Rotations	≤ 3	$O(\log n)$	Merge-based	N/A
Ordered Traversal	✓	✓	✓	✗
Range Queries	✓	✓	✓ Best	✗
Disk Friendly	✗	✗	✓ Best	✗
Search Speed	$O(\log n)$	$O(\log n)^*$	$O(\log n)$	$O(1)$ avg

✏ *AVL is faster in practice for read-heavy workloads due to shorter tree height ($\leq 1.44 \log n$ vs $\leq 2 \log n$ for RBT)

Decision Guide

- **Write-balanced in-memory:** Red-Black Tree
- **Read-heavy in-memory:** AVL Tree
- **Disk-based storage:** B+ Tree
- **Concurrent systems:** Skip List
- **Unordered lookup:** Hash Table

Learning Summary

Key Takeaways

1 Five Invariants, One Guarantee

A Red-Black Tree maintains 5 structural color invariants to keep height $\leq 2\log(n+1)$

2 Insert RED, Root BLACK

New nodes are always inserted as RED; root is always forced to BLACK after fix-up

3 Three Insert Fix-Up Cases

Uncle RED $\rightarrow$ recolor; Uncle BLACK inner $\rightarrow$ rotate+recolor $\times 2$; Uncle BLACK outer $\rightarrow$ rotate+recolor

4 Double Black for Delete

Delete introduces "Double Black" — requires up to 4 fix-up cases but at most 3 rotations total

5 Write-Heavy Champion

RBT has fewer worst-case rotations per delete than AVL $\rightarrow$ preferred for write-heavy workloads

6 Ubiquitous in Production

Powers Linux CFS, Java TreeMap/TreeSet, C++ `std::map/set`, Nginx timers, Java HashMap bins

Memory Aids / Mnemonics

"Red can't touch Red"

Property 4 in one phrase — no two consecutive red nodes on any path

"Black Height is sacred"

Property 5 — everything else can shift to maintain this invariant

"INSERT = Color Problem"

Insert worries about consecutive reds; fix by recoloring and rotating

"DELETE = Count Problem"

Delete worries about losing a black node; fix by Double Black resolution

Quick Reference & Next Steps

Revision Cheat Sheet

RED-BLACK TREE QUICK REFERENCE	
5 PROPERTIES:	
1. Every node: RED or BLACK	
2. Root: BLACK	
3. NIL leaves: BLACK	
4. RED node → both children BLACK	
5. All root→NIL paths: equal black height	
INSERT: New node = RED. Fix-up if parent = RED	
Uncle RED → Recolor (push problem up)	
Uncle BLACK, inner child → Rotate child	
Uncle BLACK, outer child → Rotate grandparent	
DELETE: If black node removed → Double Black fix-up	
Sibling RED → Rotate, convert to next	
Sibling BLACK, both nephews BLACK → Recolor, push up	
Sibling BLACK, near RED → Rotate sibling	
Sibling BLACK, far RED → Rotate parent, done	
COMPLEXITY: All operations $O(\log n)$	
INSERT ≤ 2 rotations DELETE ≤ 3 rotations	
HEIGHT ≤ $2 \times \log_2(n+1)$	

Common Interview Questions

1. List all 5 Red-Black Tree properties
2. Why are new nodes always inserted as RED?
3. What is black height, and why must it be equal on all paths?
4. How many rotations can an insert trigger? A delete?
5. How is a Red-Black Tree related to a 2-3-4 B-Tree?
6. Why does Java use Red-Black Trees in TreeMap instead of AVL Trees?
7. What is the "Double Black" concept in deletion?
8. How does the Linux kernel save memory by packing the color bit into the parent pointer?
9. When would you choose AVL over Red-Black, and vice versa?
10. What happens in Java's HashMap when a bucket exceeds 8 entries?

Recommended Next Topics

B-Trees

Multiway generalization; disk-based analog

B+ Trees

Used in every relational database index

Interval Trees

RBT augmented for interval overlap queries

Order-Statistic Trees

RBT augmented for rank/select

Skip Lists

Probabilistic alternative; easier concurrent implementation

Concurrent Data Structures

Lock-free and fine-grained concurrent tree variants